

INFLUXDB

10
HOUR
LIMIT

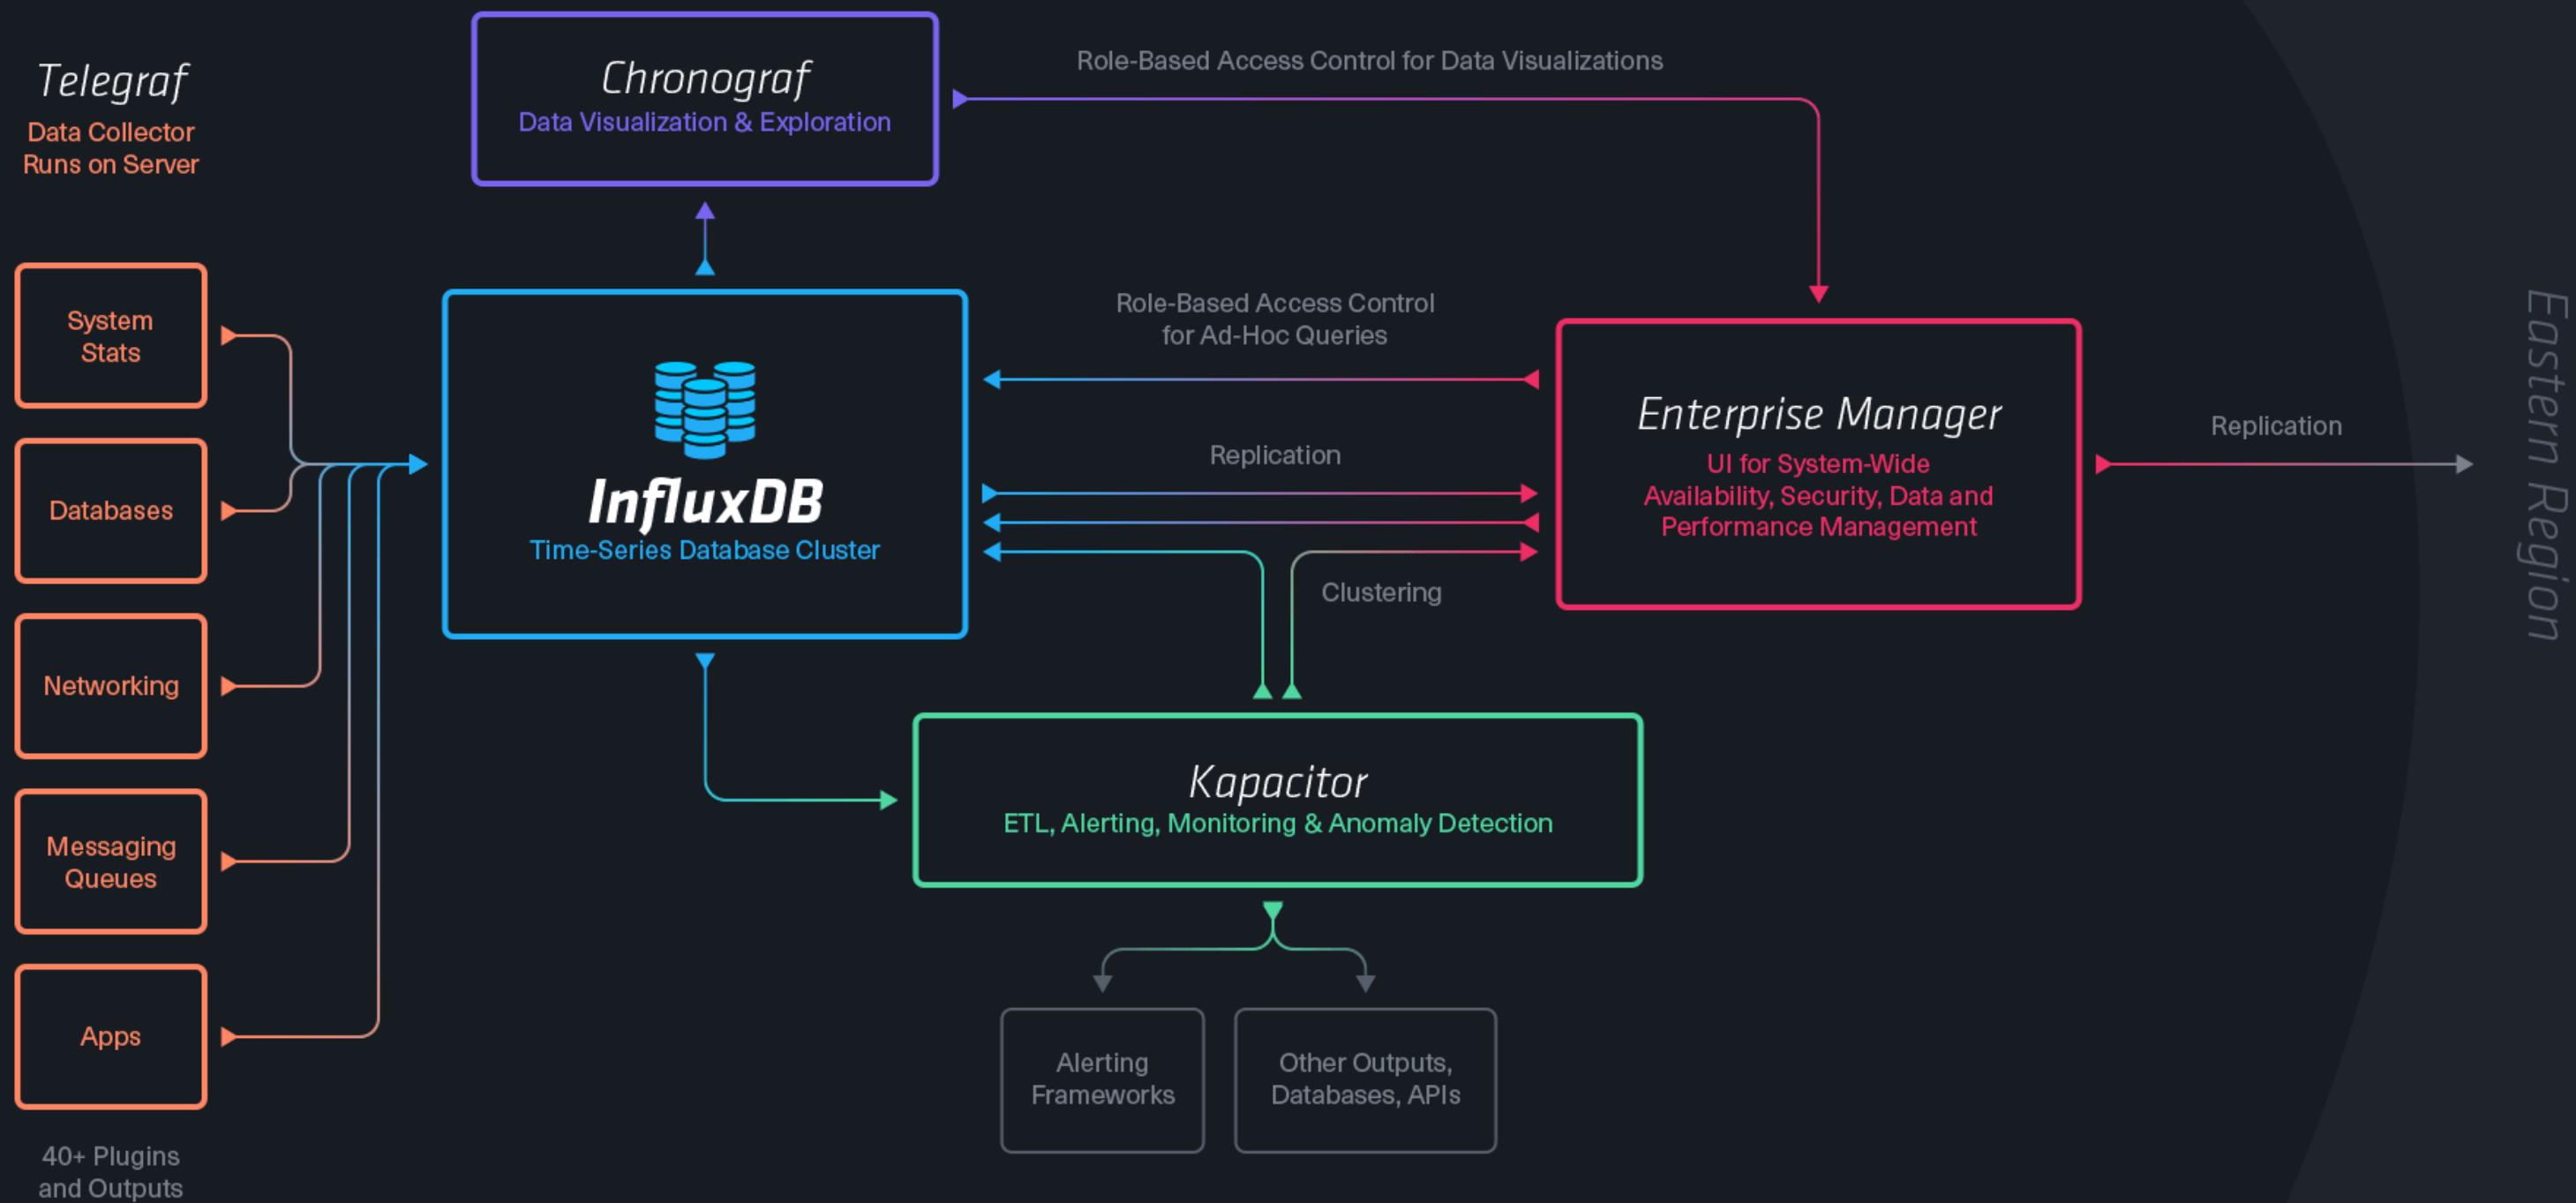
10
HOUR
LIMIT

073N0170

073N0162

TIME SERIES DATABASE

Western Region



USE CASES

DATA MODEL

- ▶ String
- ▶ Integer
- ▶ Floats
- ▶ Boolean

```
influxd -config /usr/local/etc/influxdb.conf
```

Query: select mean(value) from cpu where time >= '2015-12-29 10:04:40' AND time < '2015-12-29 10:04:59.0' group by time(10s)

Query Templates ▾

cpu

time	mean
2015-12-29T10:04:40Z	0.5466666666666666
2015-12-29T10:04:50Z	1

QUERY LANGUAGE

```
SELECT
  COUNT(duration) as count_duration,
  MIN(duration) as min,
  MAX(duration) as max,
  MEAN(duration) as MEAN
FROM events
WHERE time > now() - 1h
GROUP BY time(30s)
```

```
INSERT cpu_load,server_name=gilbert value=2
```

```
curl -i -XPOST 'http://localhost:8086/write?db=mydb'  
--data-binary 'cpu_load,server_name=gilbert value=2'
```

RETENTION

```
CREATE RETENTION POLICY two_hours  
ON food_data DURATION 2h REPLICATION 1 DEFAULT
```

DOWNSAMPLING

```
CREATE CONTINUOUS QUERY cq_30m ON food_data
BEGIN
  SELECT mean(temperature) AS mean_temperature,
         mean(delivery_time) AS mean_delivery_time
  INTO food_data."default".downsampled_orders
  FROM orders
  GROUP BY time(30m)
END
```

TELEGRAF

INPUT PLUGINS

aerospike, apache, bcache, couchdb, disque, dns query time, docker, dovecot, elasticsearch, exec, haproxy, httpjson, influxdb, jolokia, leoofs, lustre2, mailchimp, memcached, mesos, mongodb, mysql, net_response, nginx, nsq, phpfm, phusion passenger, ping, postgresql, powerdns, procstat, prometheus, puppetagent, rabbitmq, raindrops, redis, rethinkdb, riak, sensors, snmp, sql server, twemproxy, zfs, zookeeper, win_perf_counters, system, cpu, mem, net, netstat, disk, diskio, swap, statsd, mqtt_consumer, kafka_consumer, nats_consumer, github_webhooks

OUTPUT PLUGINS

influxdb, amon, amqp, aws kinesis,
aws cloudwatch, datadog, graphite,
kafka, librato, mqtt, nsq, opensdsb,
prometheus, riemann

```
telegraf -sample-config  
-input-filter cpu  
-output-filter influxdb  
> telegraf.conf
```

```
[[outputs.influxdb]]  
urls = ["http://localhost:8086"]  
database = "telegraf"  
precision = "s"
```

```
[[inputs.cpu]]  
percpu = true  
totalcpu = true  
drop = ["time_*"]
```

telegraf -config telegraf.conf

CHRONOGRAF

chronograf



CANCEL

Add New Server

InfluxDB Servers

Manage which servers you'd like to use as data sources for Chronograf.

Add new server

NICKNAME

InfluxDB-1

HOST

localhost

PORT

8086

 Use SSL

HTTP BASIC AUTH USERNAME (OPTIONAL)

HTTP BASIC AUTH PASSWORD



Validate server before add

Add

< CANCEL

Average idle CPU usage

DONE >

Auto Refresh: None ▾

🕒 Past 15 minutes ▾

0

BUILDER

SELECT value FROM "measurement" WHERE tmpltime()

FILTER BY

Servers: InfluxDB-1

Databases: telegraf

Retention Policies: default

+ ADD QUERY

Make Default

EXTRACT BY

value BY Select a function

Select a field BY

GROUP BY

Apply

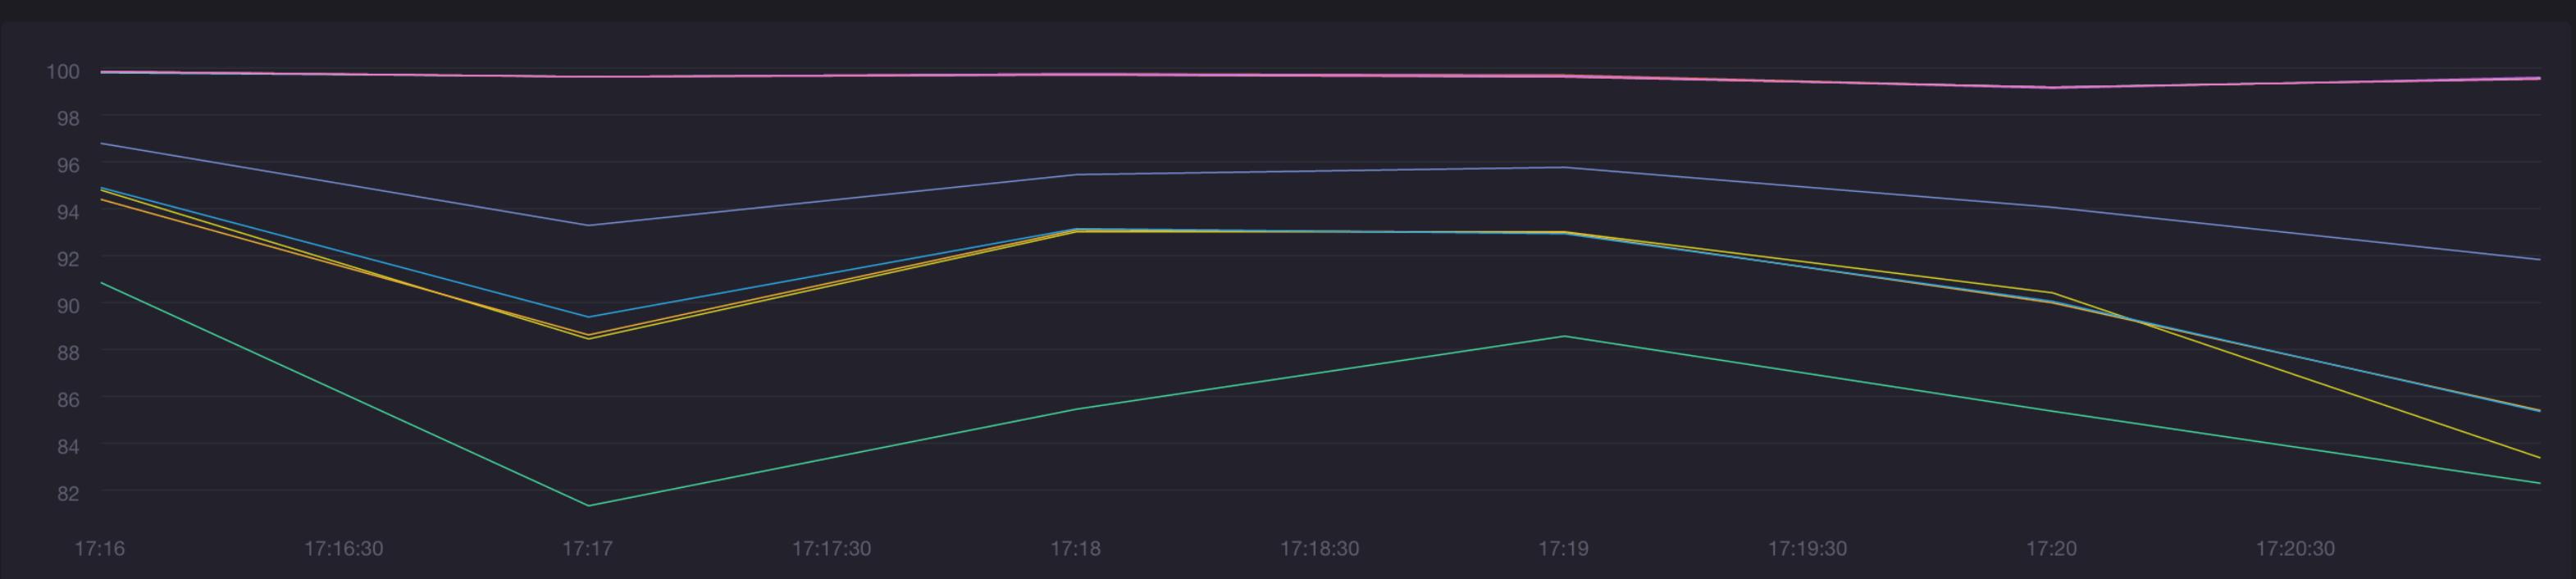


< CANCEL

Average idle CPU usage

DONE >

Auto Refresh: None

Use WHERE tmptime()
in a query to enable time
range controls.

BUILDER



SELECT mean("usage_idle") FROM "telegraf"."default"."cpu" WHERE time > now() - 5m GROUP BY time(1m), "cpu"



FILTER BY

cpu

6

Select a tag key

EXTRACT BY

usage_idle

BY mean

Select a field

BY Select a function

GROUP BY

+ ADD QUERY

cpu

1m

Select a Tag

Add Visualization

Add From Existing Visualizations

Name this Visualization

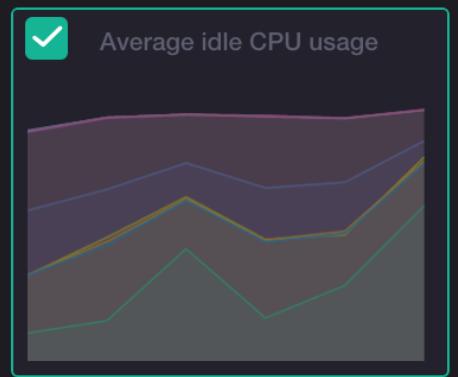
Cancel

Save

< CANCEL

Select Visualizations

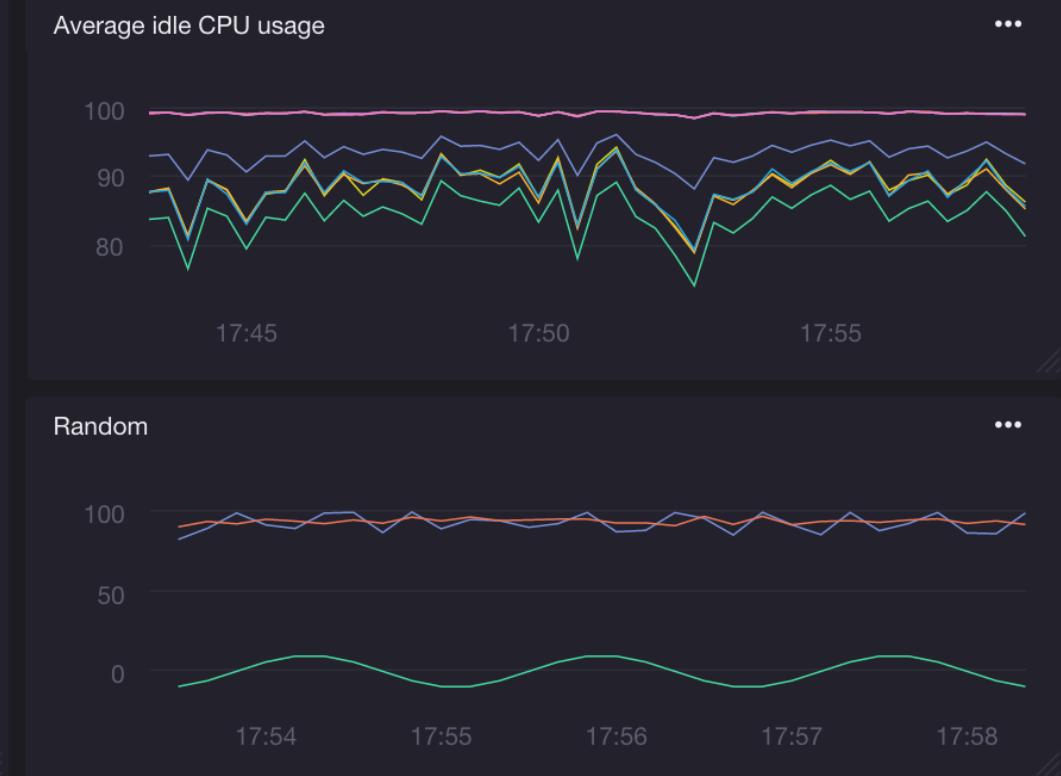
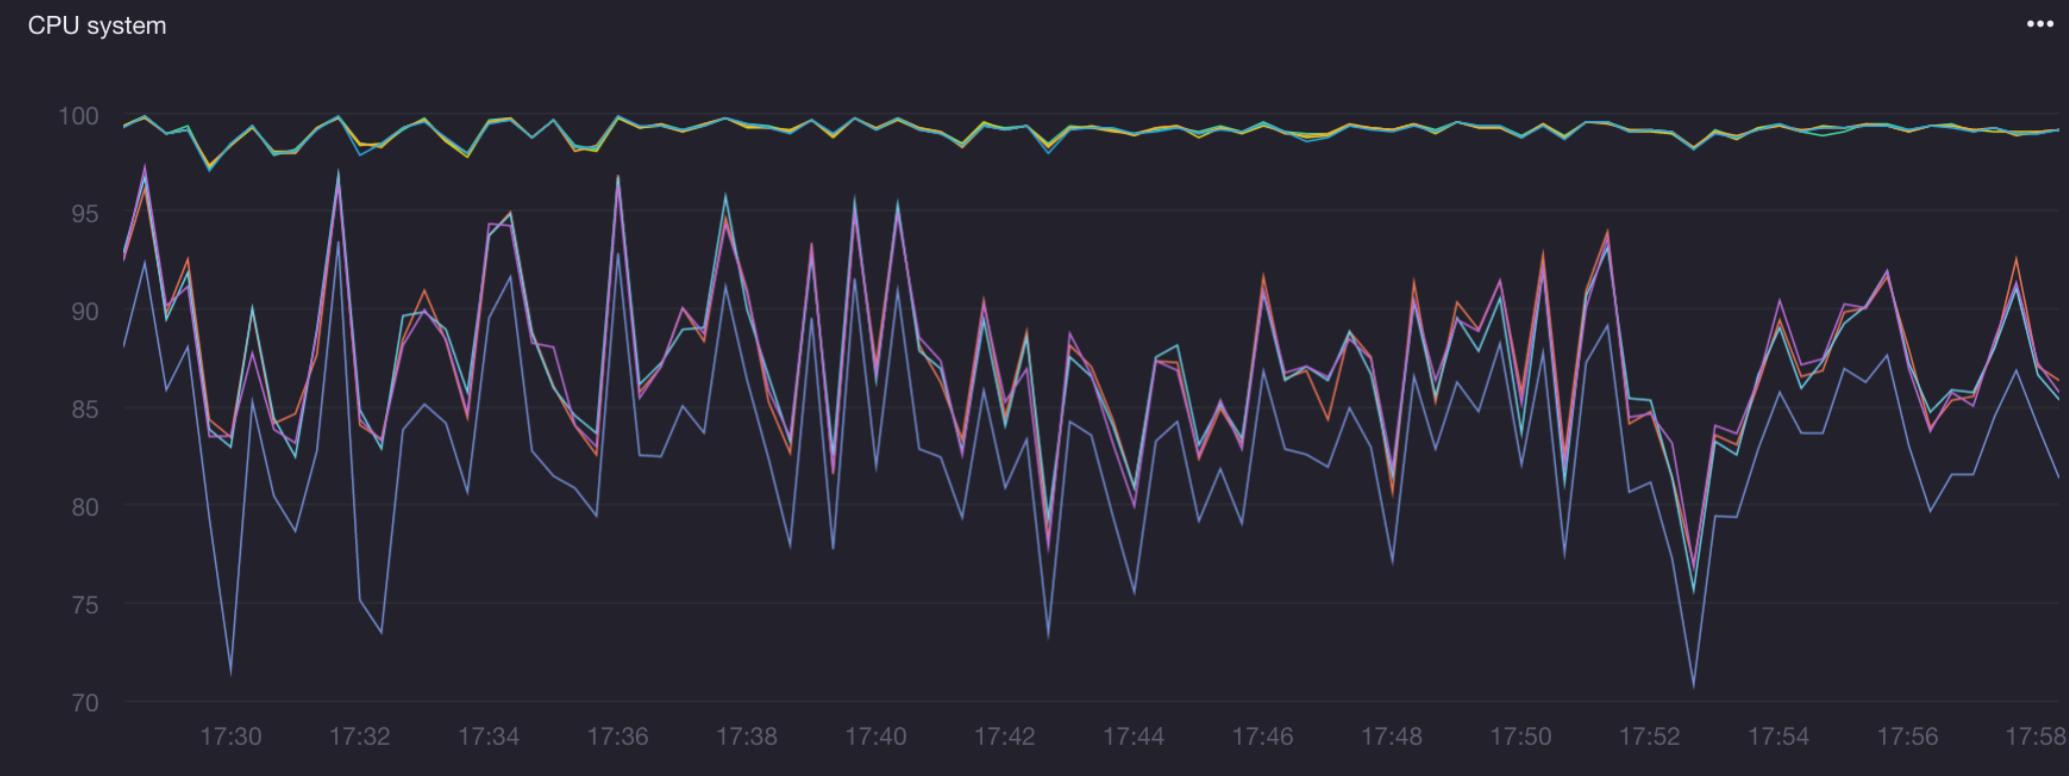
Add Visualizations to Dash



[Add Visualization](#)

Auto Refresh: None

Past 5 minutes



*GitLab Performance Monitoring: a
monitoring system using InfluxDB [...]
then visualized using Grafana*

- [Gitlab Performance Blogpost](#)

KAPACITOR

kapacitord config > kapacitor.conf

kapacitord -config kapacitor.conf

CPU_ALERT.TICK

stream

```
// Select just the cpu measurement
.from().measurement('cpu')
.alert()
.crit(lambda: "usage_idle" < 70)
// Whenever we get an alert write it to a file.
.log('/tmp/alerts.log')
```

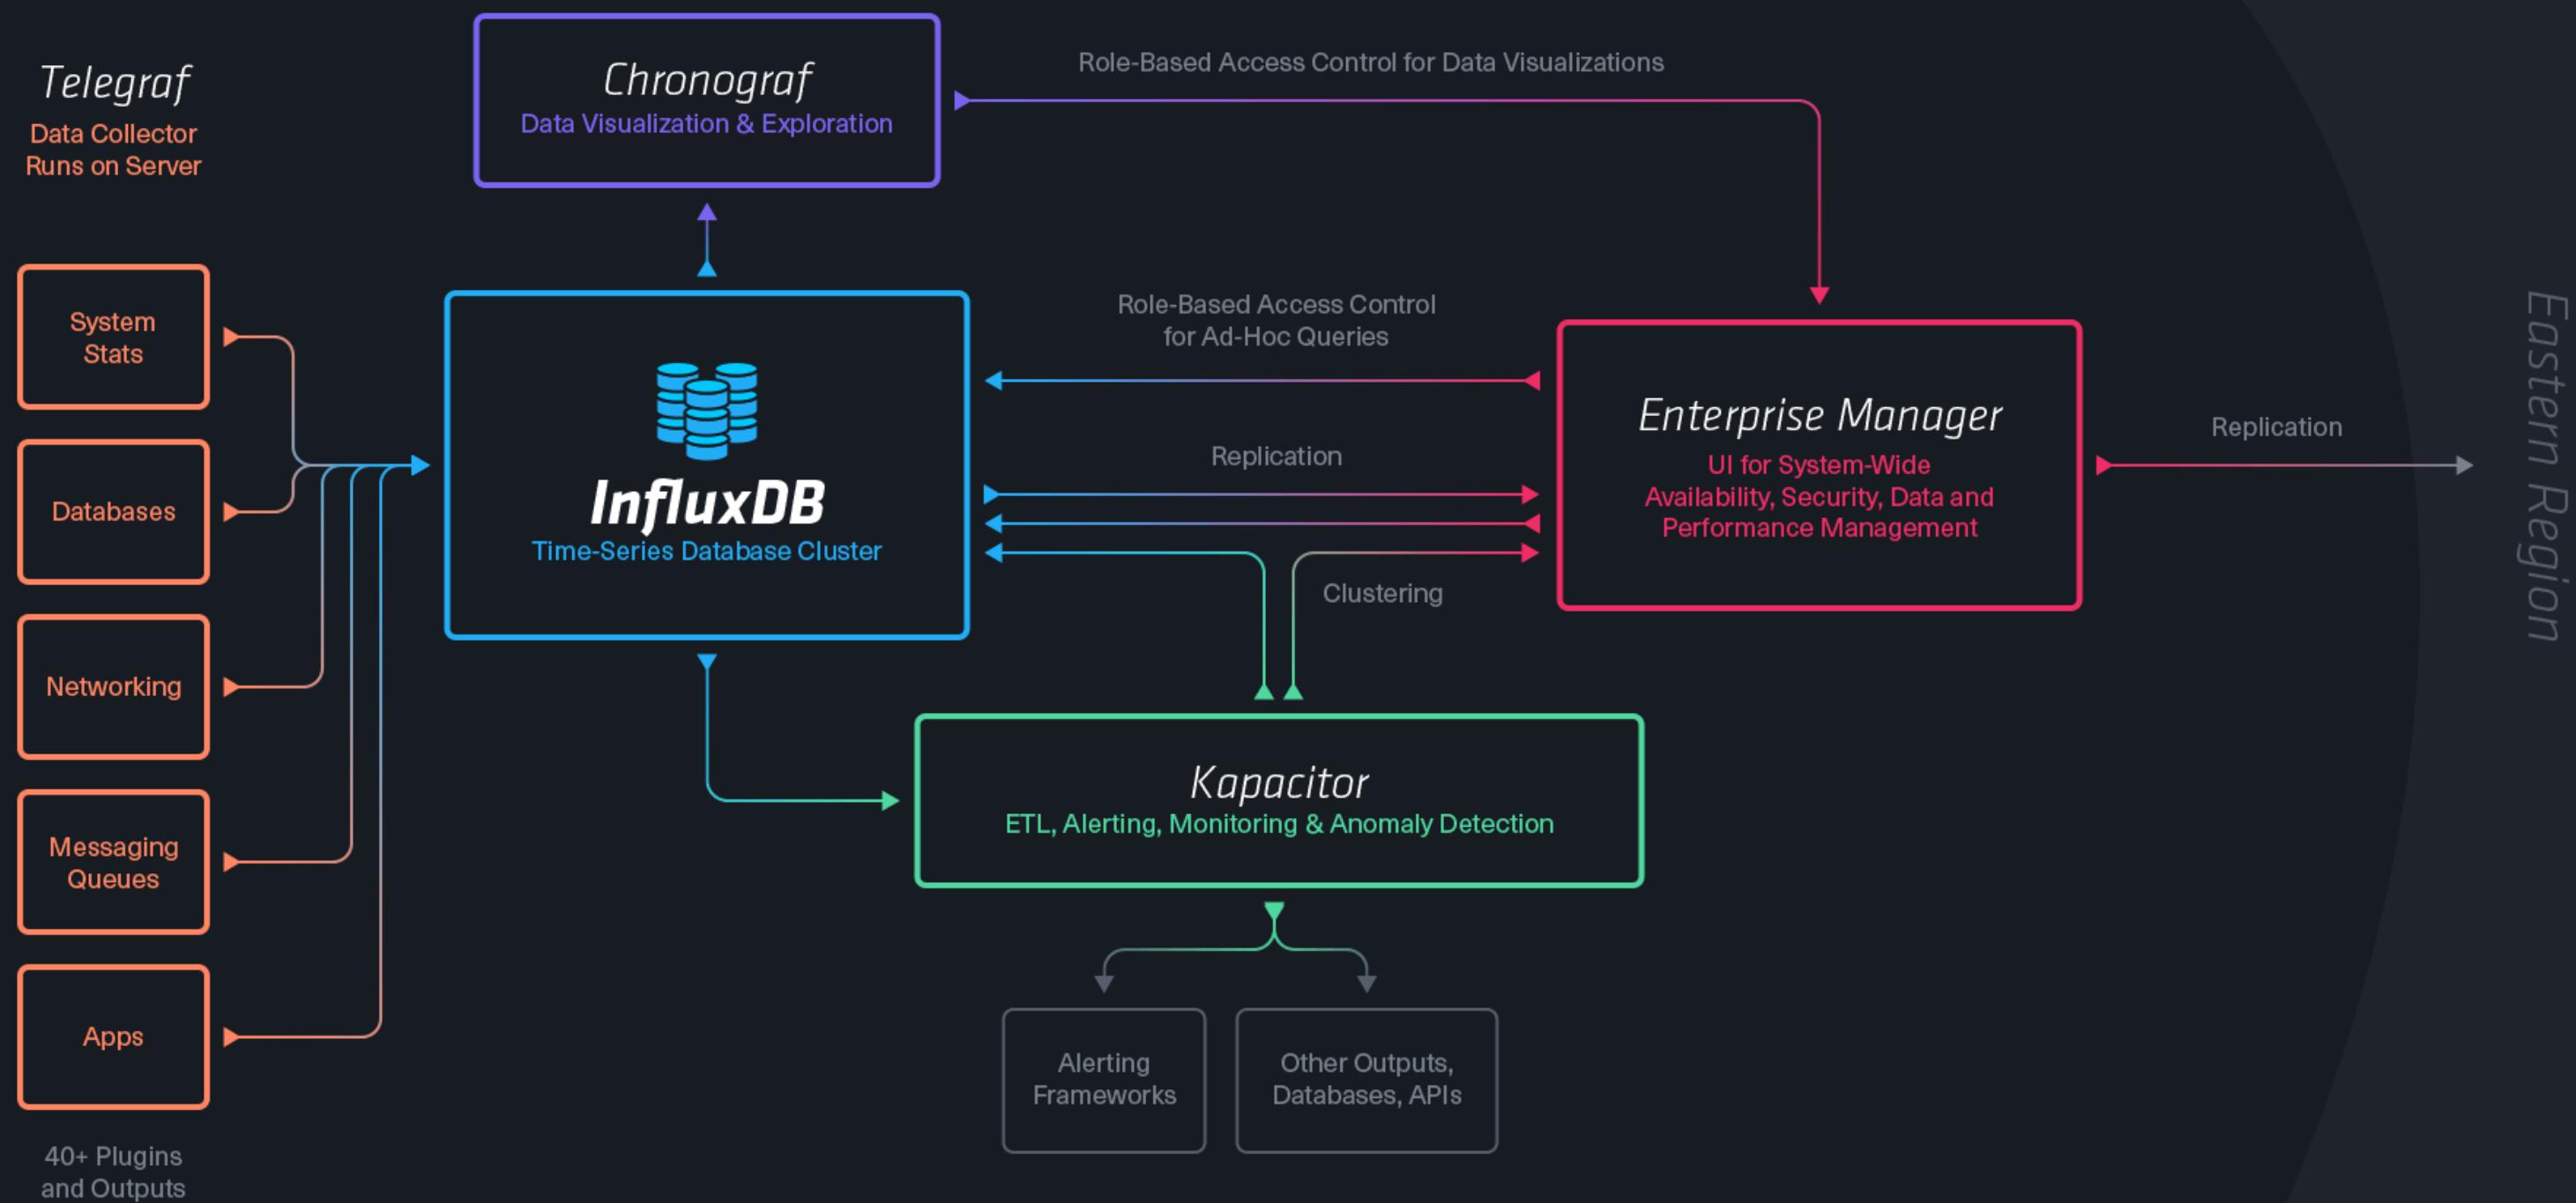
```
kapacitor define \
  -name cpu_alert \
  -type stream \
  -tick cpu_alert.tick \
  -dbrp kapacitor_example.default
```

```
kapacitor enable cpu_alert
```

```
stream
  .from().measurement('cpu')
  .alert()
    // Compare values to running mean and standard deviation
    .crit(lambda: sigma("usage_idle") > 3)
    .log('/tmp/alerts.log')
```

```
stream
  .from().measurement('cpu_usage_idle')
  .groupBy('host')
  .window().period(1m).every(1m)
  .mapReduce(influxql.mean('value'))
  .eval(lambda: 100.0 - "mean").as('used')
  .alert()
    .message('{{ .Level}}: {{ .Name }}/{{ index .Tags "host" }} has ' +
              'high cpu usage: {{ index .Fields "used" }}')
    .warn(lambda: "used" > 70.0)
    .crit(lambda: "used" > 85.0)
    .slack().channel('#alerts')
    .pagerDuty()
```

Western Region



Bodo Tasche
@bitboxer

- CTO bitcrowd

CREDITS

- * Time Series Image CC-BY-2.0 Ian Sane
<https://www.flickr.com/photos/31246066@N04/5261957053>